View, add and edit your notes in the app

# Coursera - Creating an AWS EC2 Autoscaling Group using Load Balancer

Generated on December 20, 2023

## Summary

| Notes | Screenshots | Bookmarks |
| --- | --- | --- |
| 22 | 25 | 0 |

Task 7:

Load Testing Our Auto Scaling Group

0:08

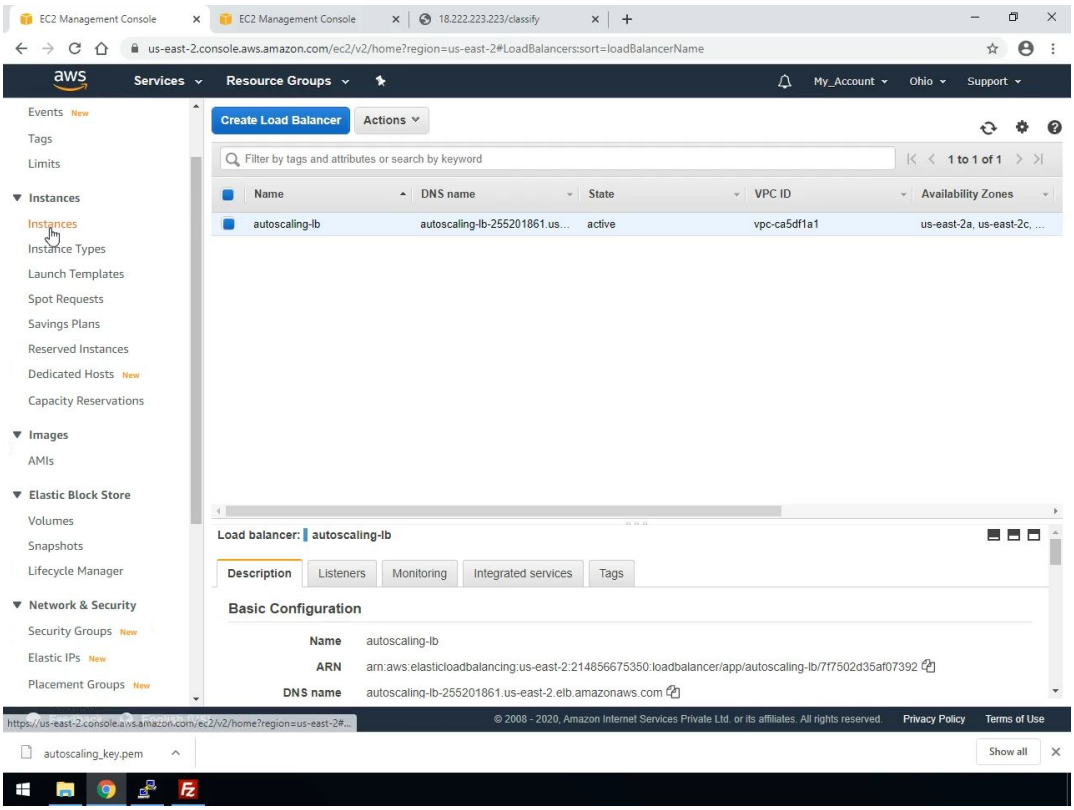


0:31

the last instance is in auto-scaling group

so we will terminate this and see what happens.

0:51

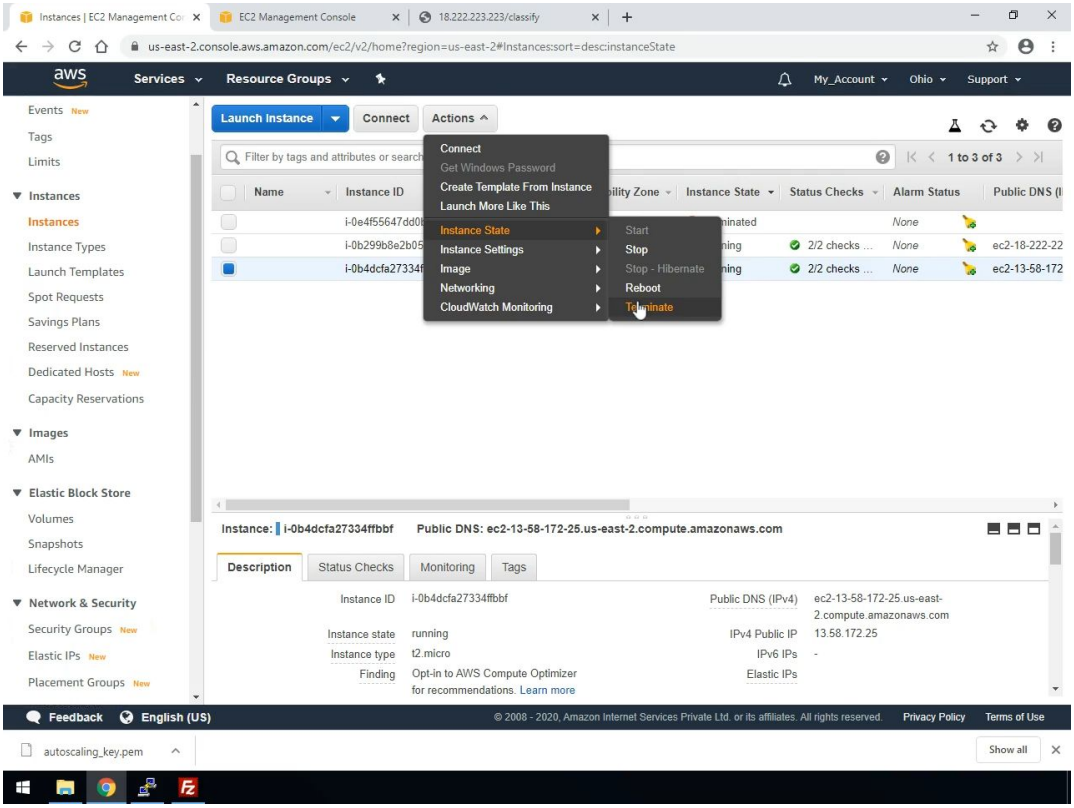


1:12

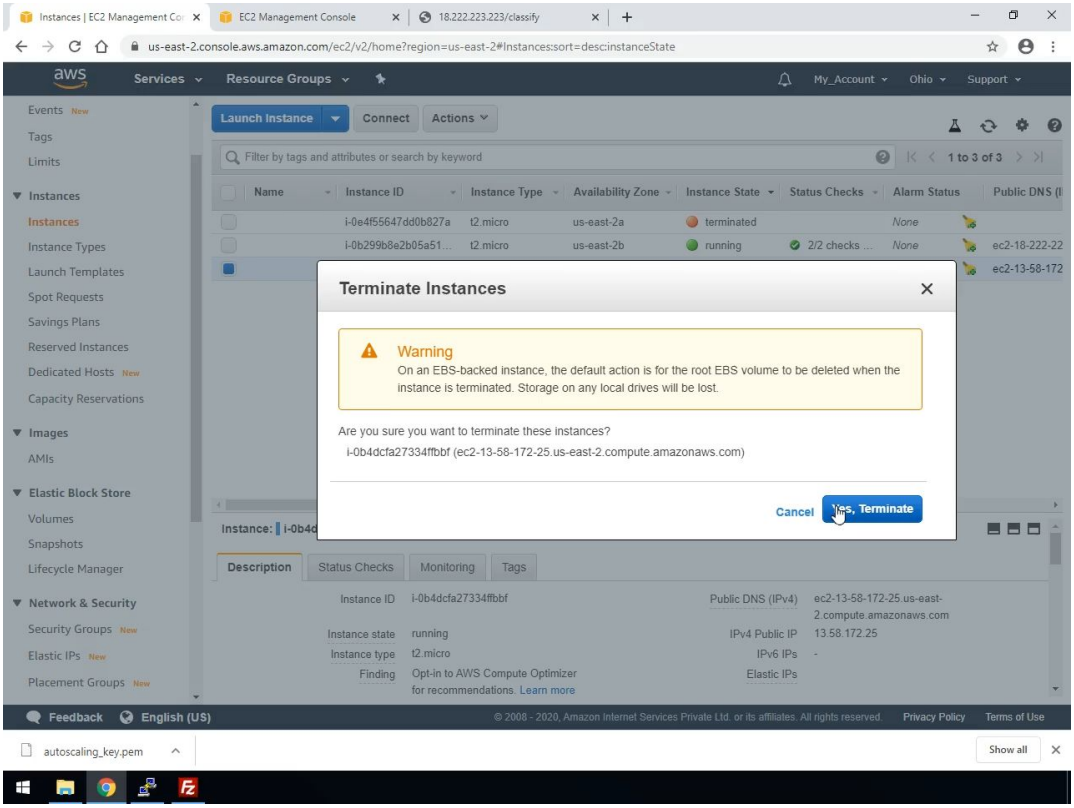


1:16

now Our Load Balancer Would not Work,

because we do ot have any instance.

1:29

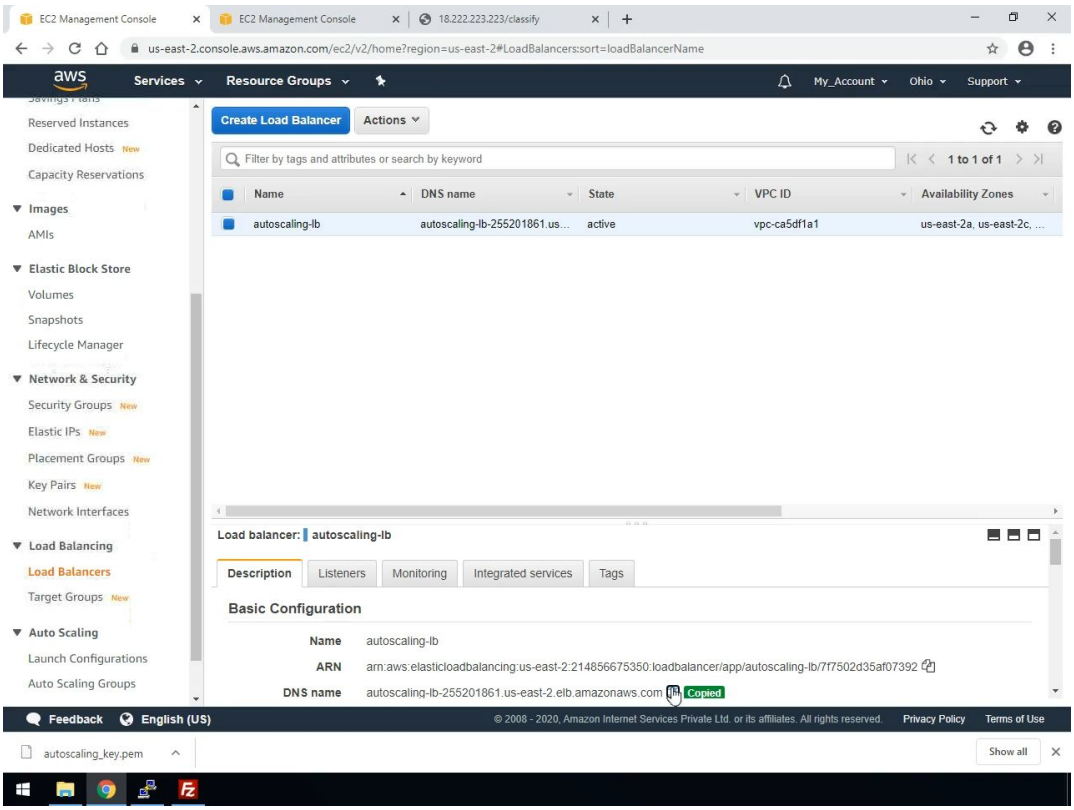


1:36

copy the DNS name

1:36

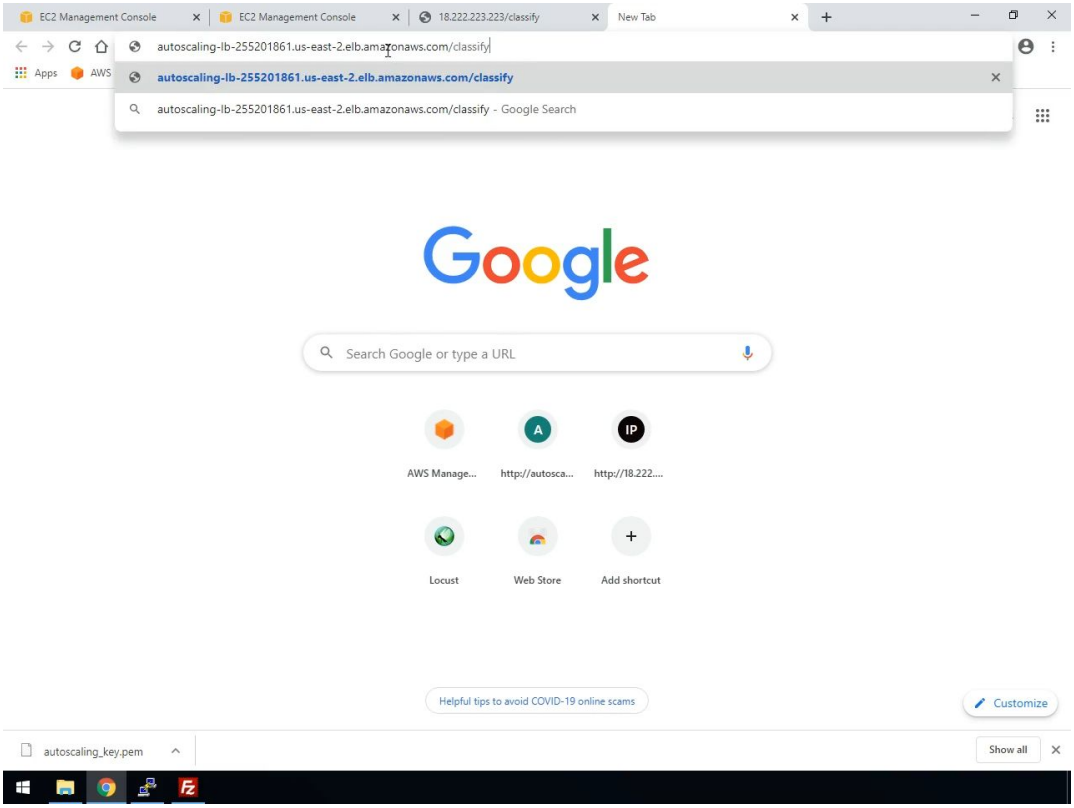

1:47


503 Service Temporarily Unavailable


1:49

let’s wait for a minute or two,

the load balancer will do couple of health checks and

it will se if the instance is returning 200 status code

2:13

so here it is not returning OK status

as we see it is returning the 500 error code

2:30

So our auto scaling group will start a new instance

2:34

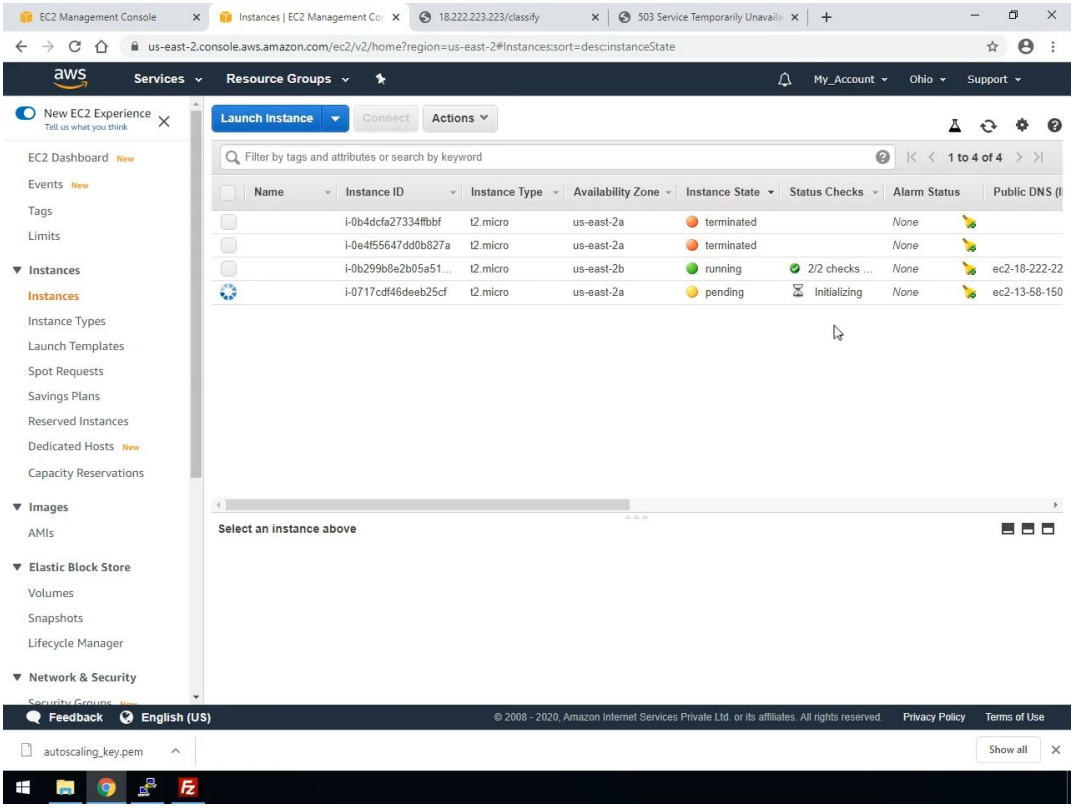


2:47

auto scaling group has started a new instance

2:52

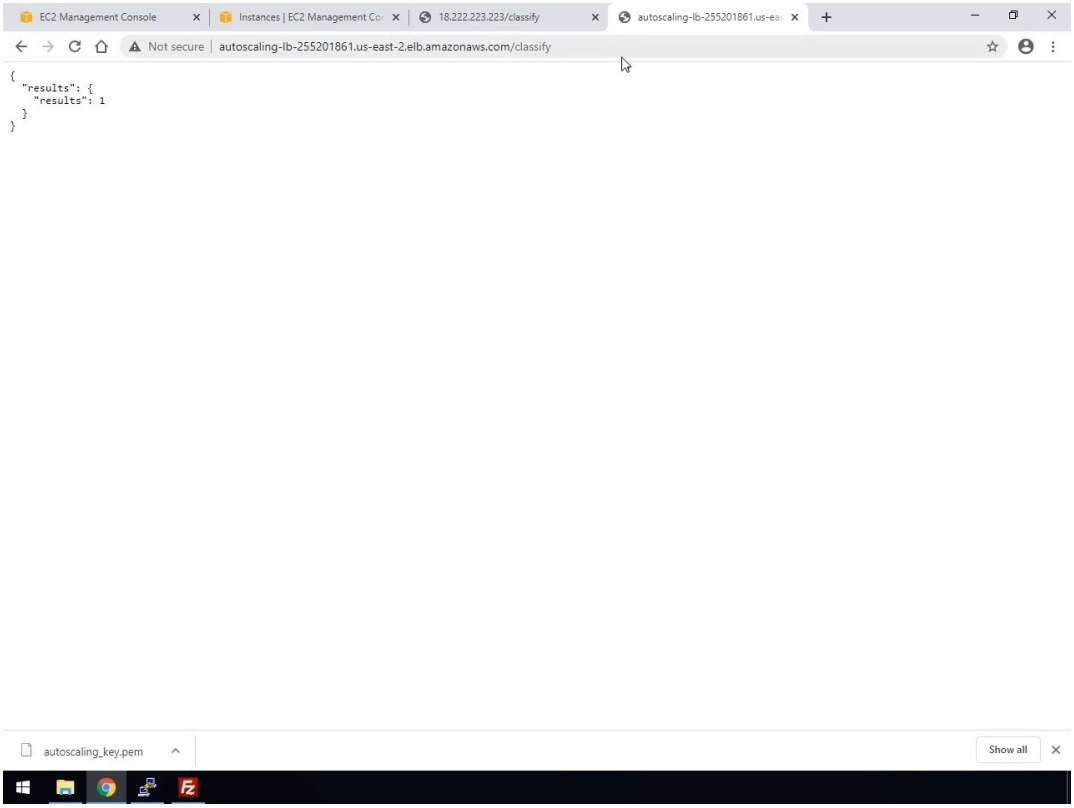

3:22

now it’s working perfectly!!

3:24

Now our final task is to laod test our application

3:32

we will use LOCUST.(search web)

so basically we will be simulating a lot of requests to our application, which increase its CPU utilixation

4:02

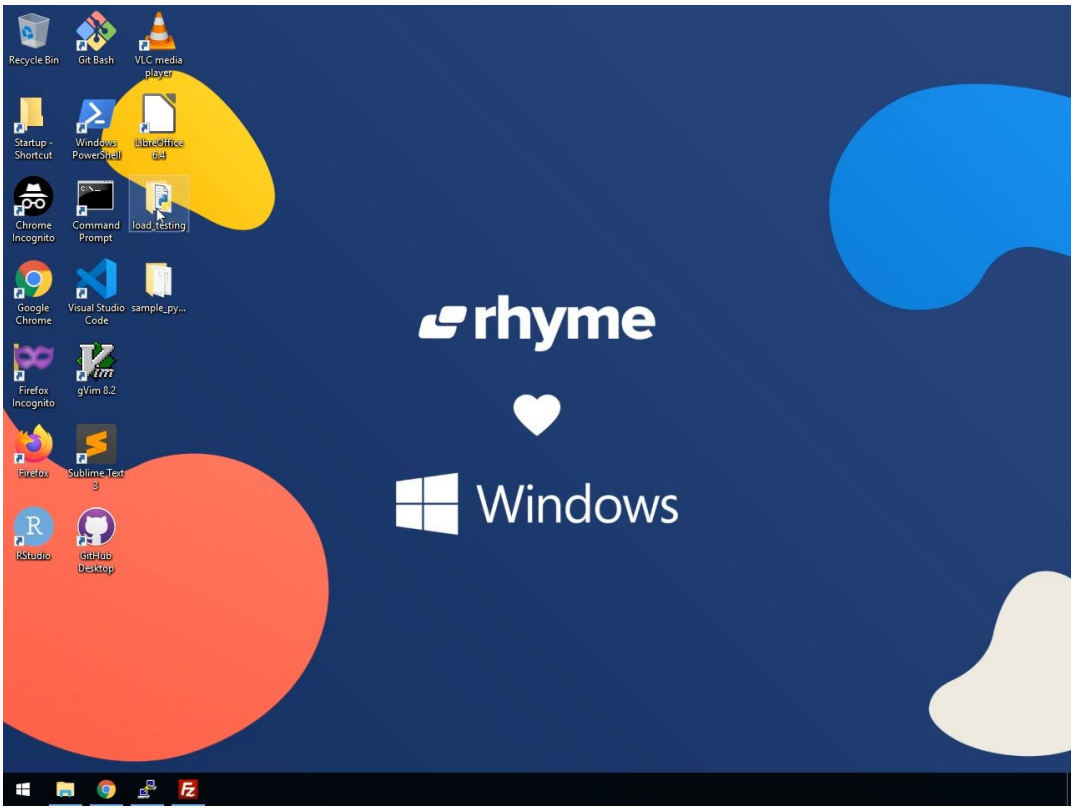


4:20

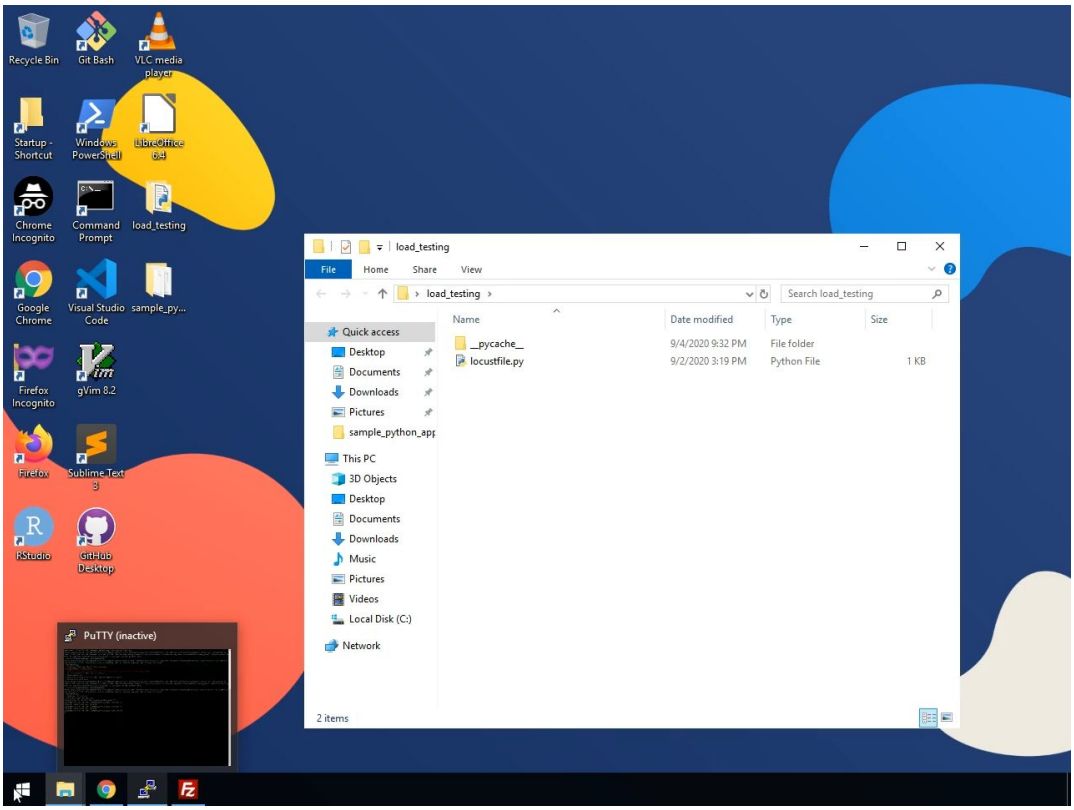


4:23

do not write classify in the last here!!!

5:10

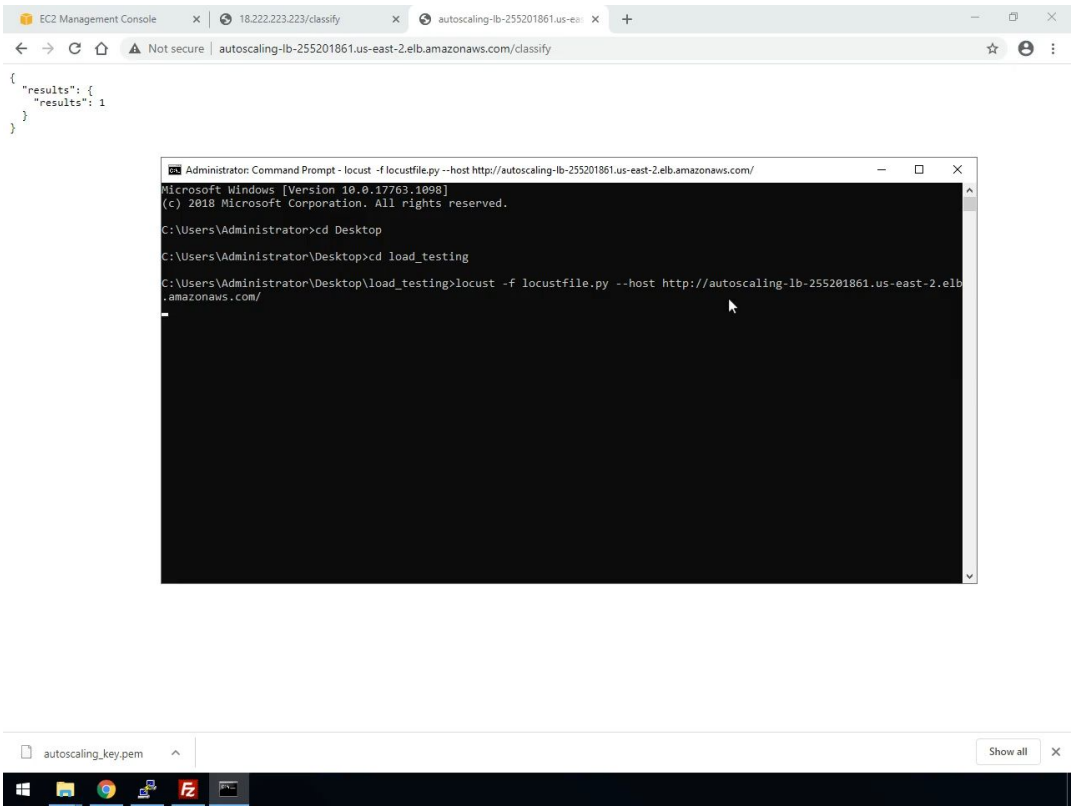


5:10

so our locust application should be running on 8089 port.

5:21

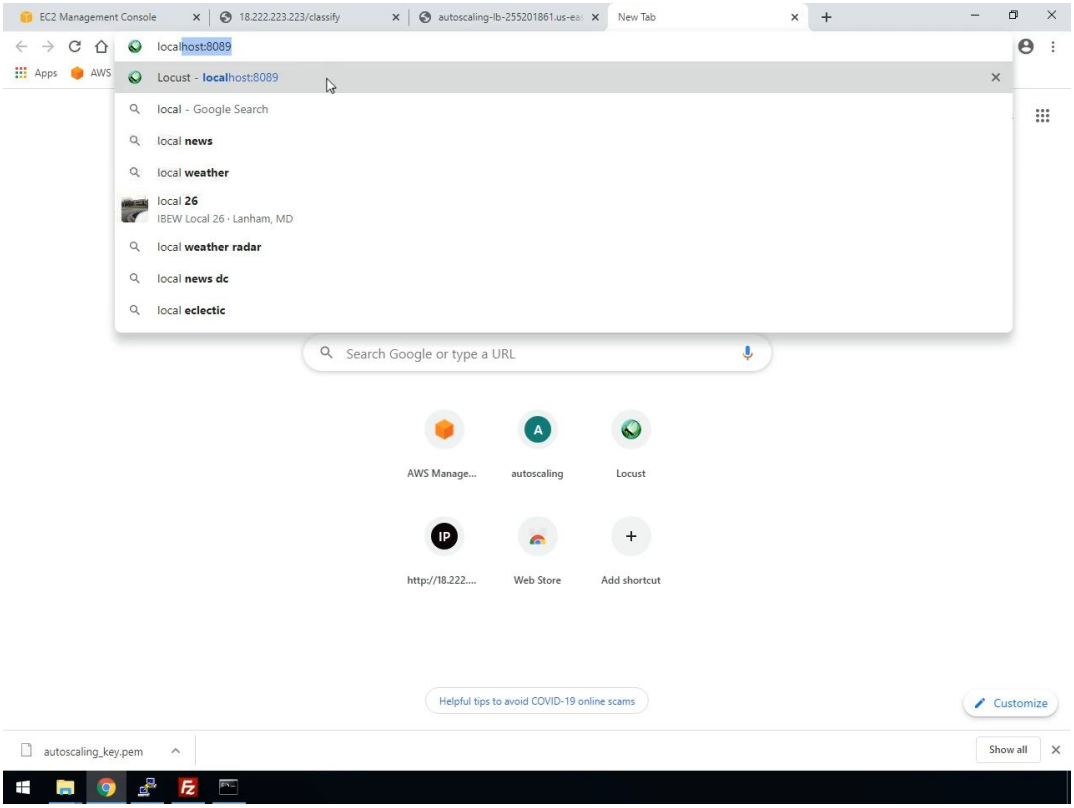


5:28

here we can see our auto-scaling group address

5:34

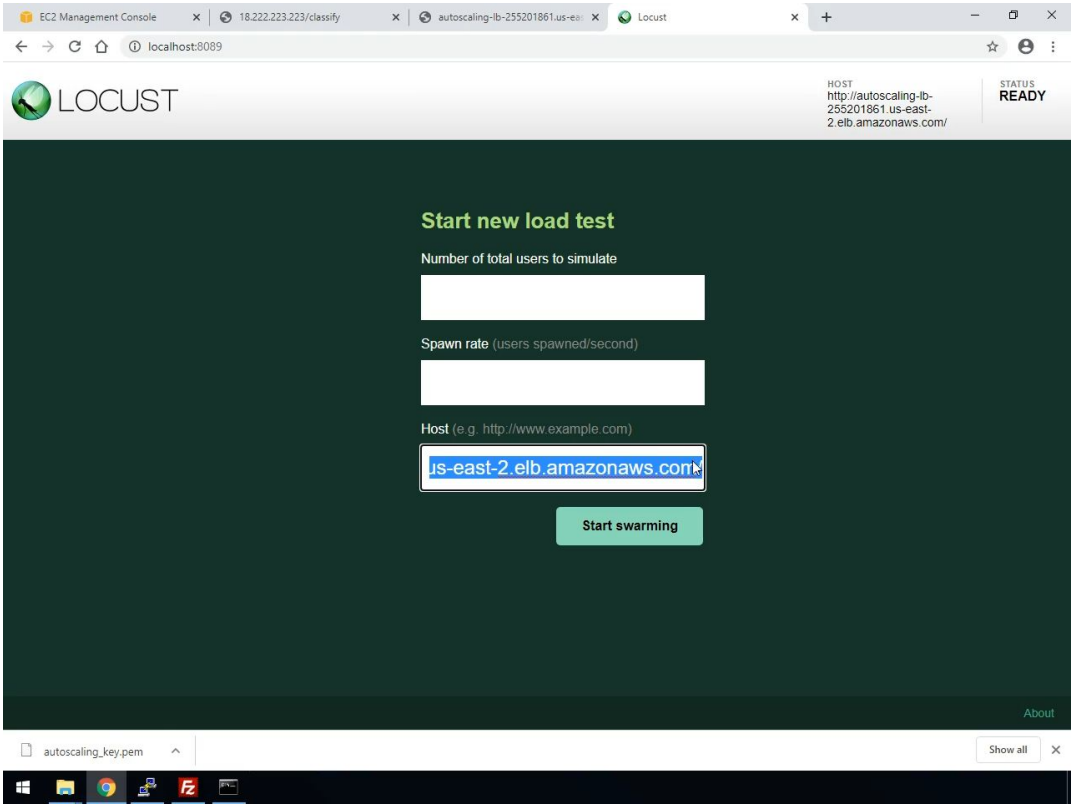

5:34

we can specify the number of users we want to stimulate

in our host >> Avg.CPU utilization is >> 50%

if above 50% it will launch a new insatnce

5:44

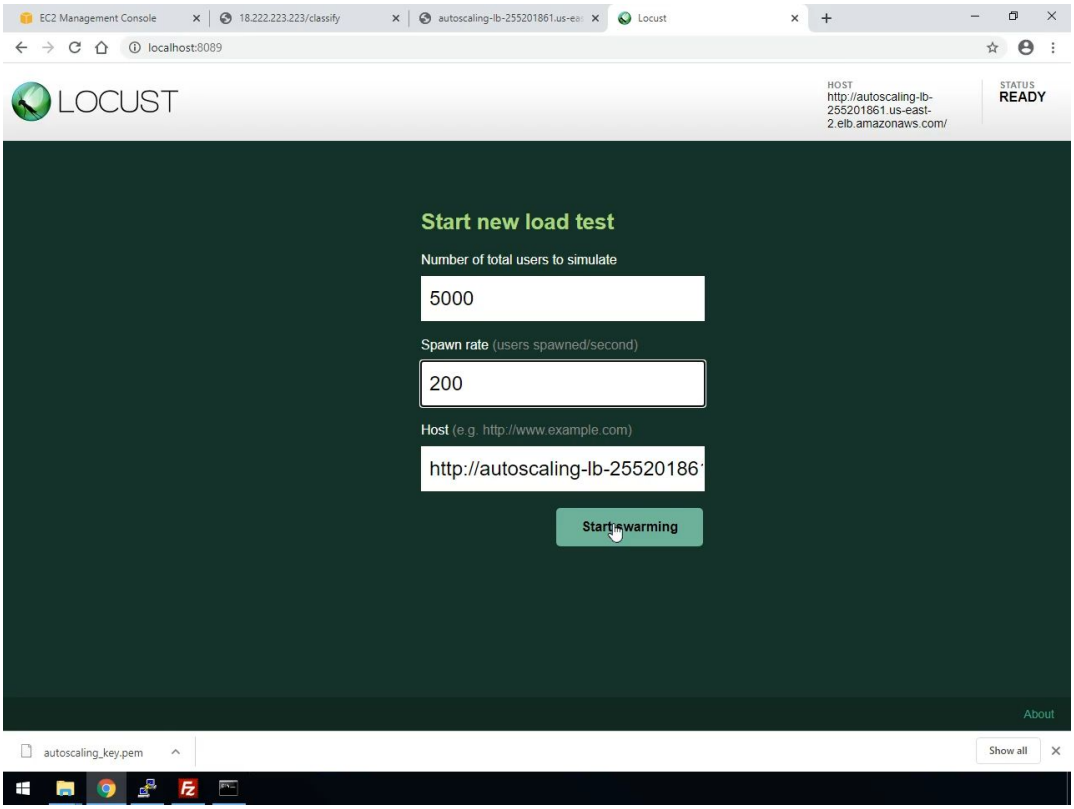


6:01

it will send 500 request per second (RPS)

6:42

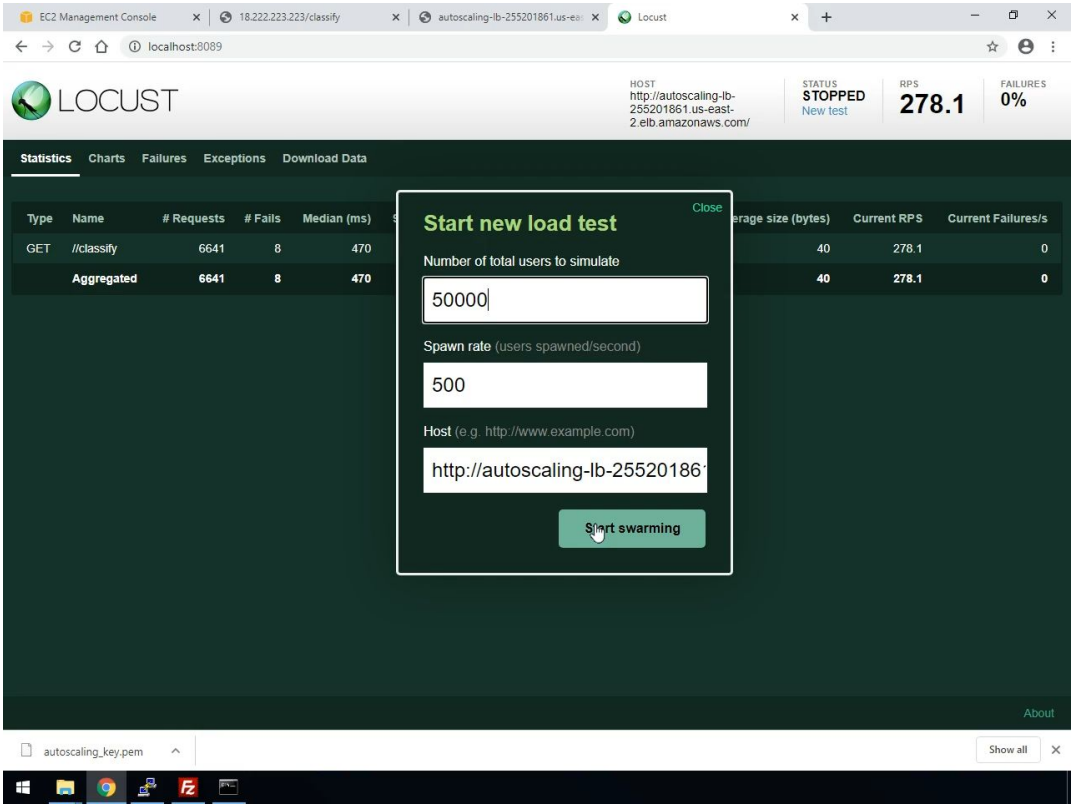


6:42

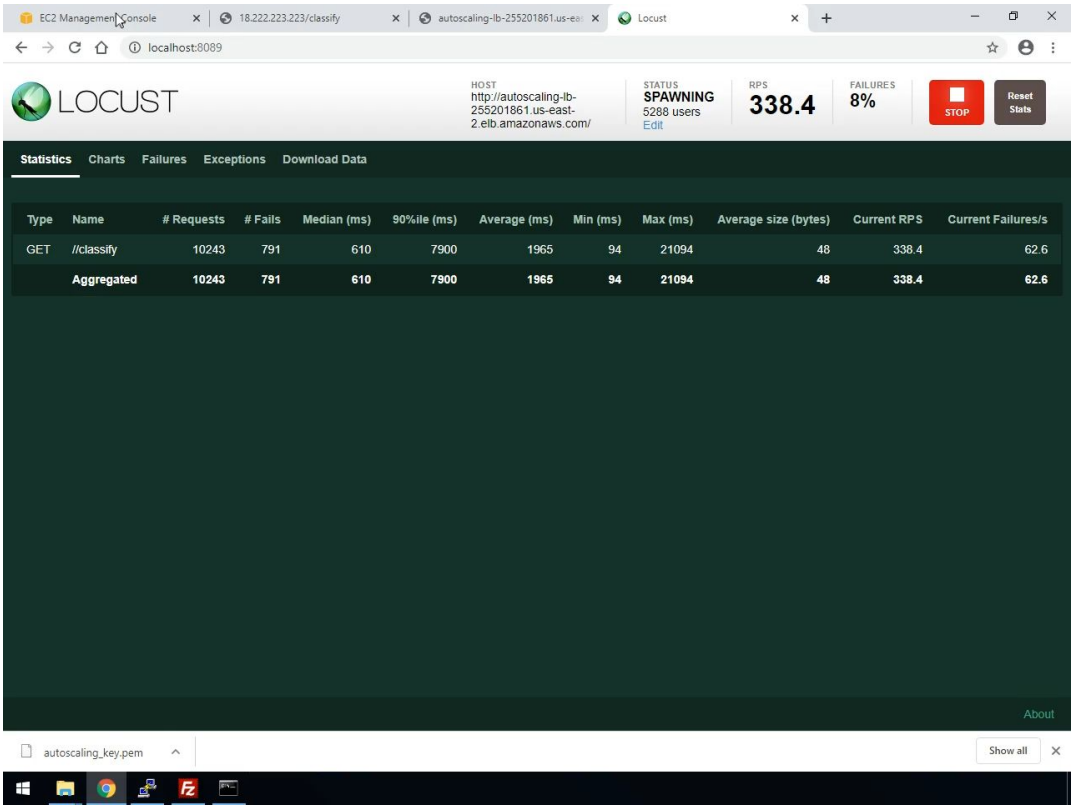


7:19

request failing start

CPU Utilixation is up from the limit

so let’s check

7:19

it will take time

wait and then refresh the instances tab

and you will see auto-scaling group will launch a new instance

8:40

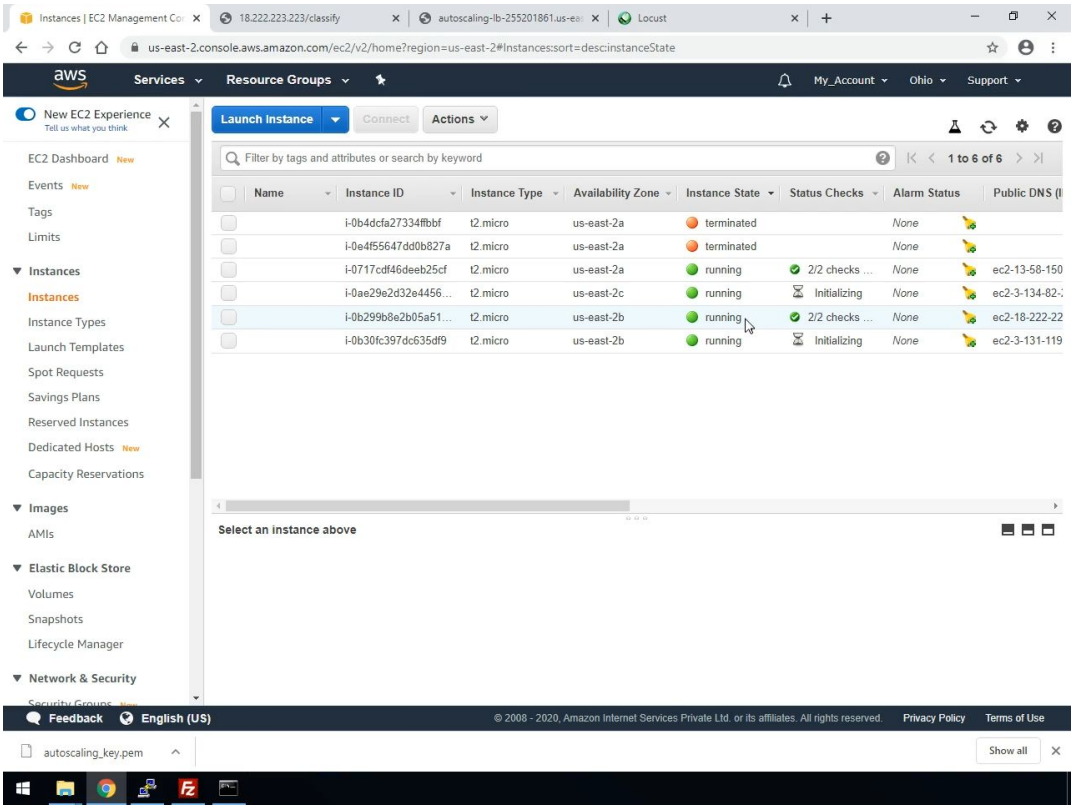


9:03

two instance initialize because there have lot of requests send to our ec2-insatnce

9:10

so when these instance start running it will decrease the fail requests

9:40

## Deleting Everything >> Prevent Billing Cost

10:40

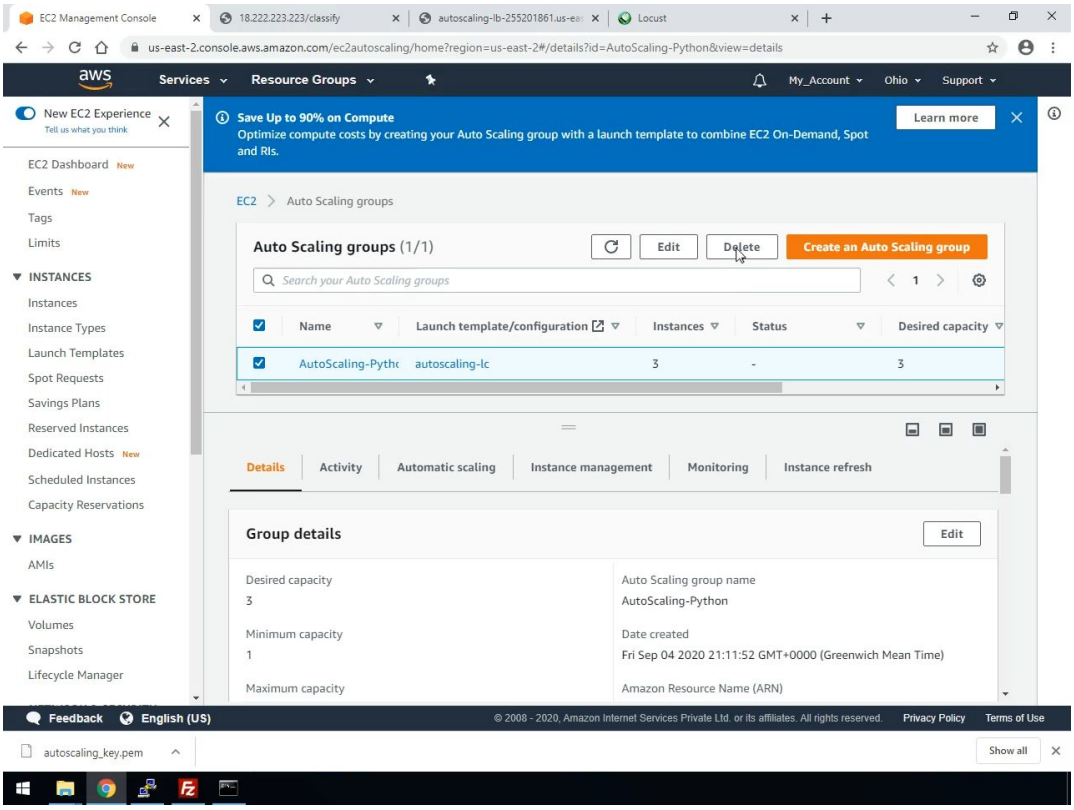


11:16

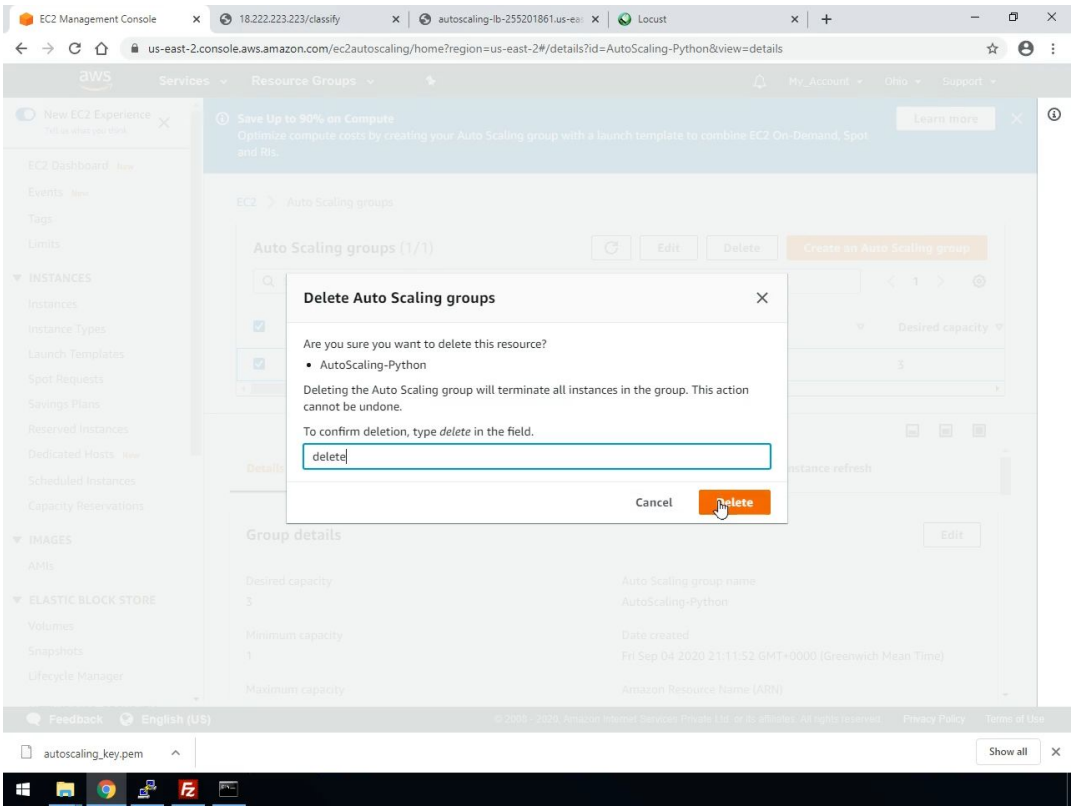


11:24

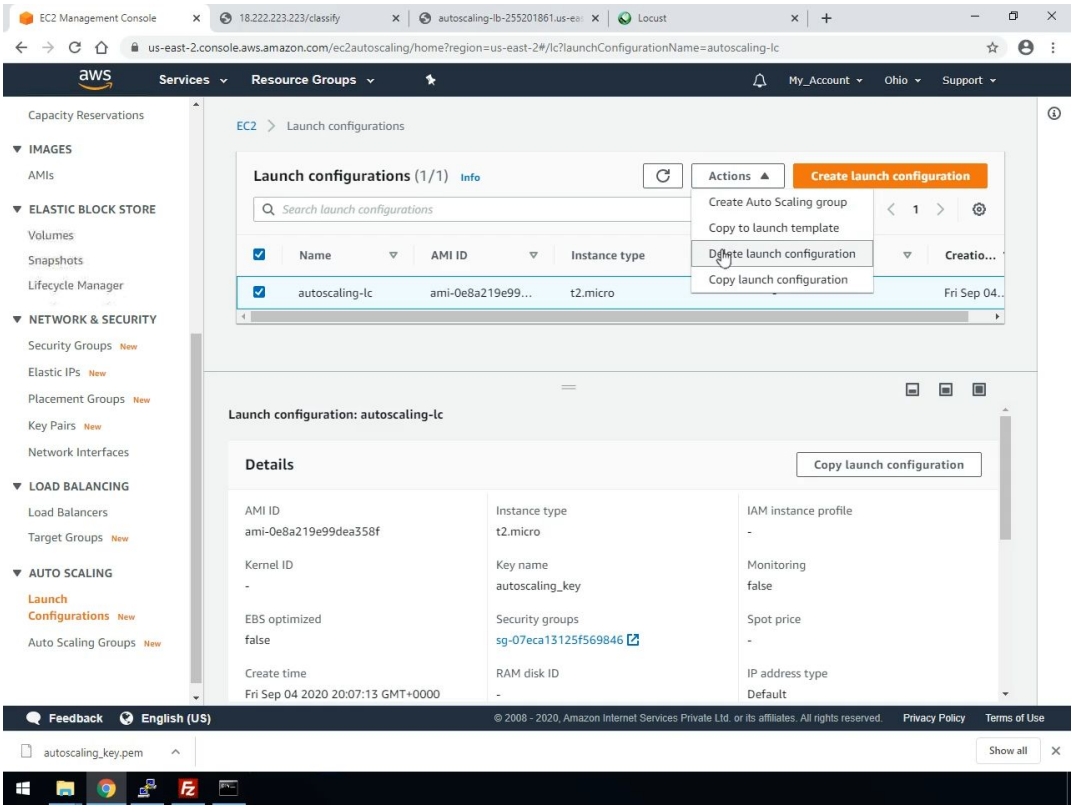


11:42

1-delete auto scaling group

2-delete launch configurations

3-delete laod balancer

4-delete Target Groups

5-delete ec2-instances

6-delete AMI

11:48

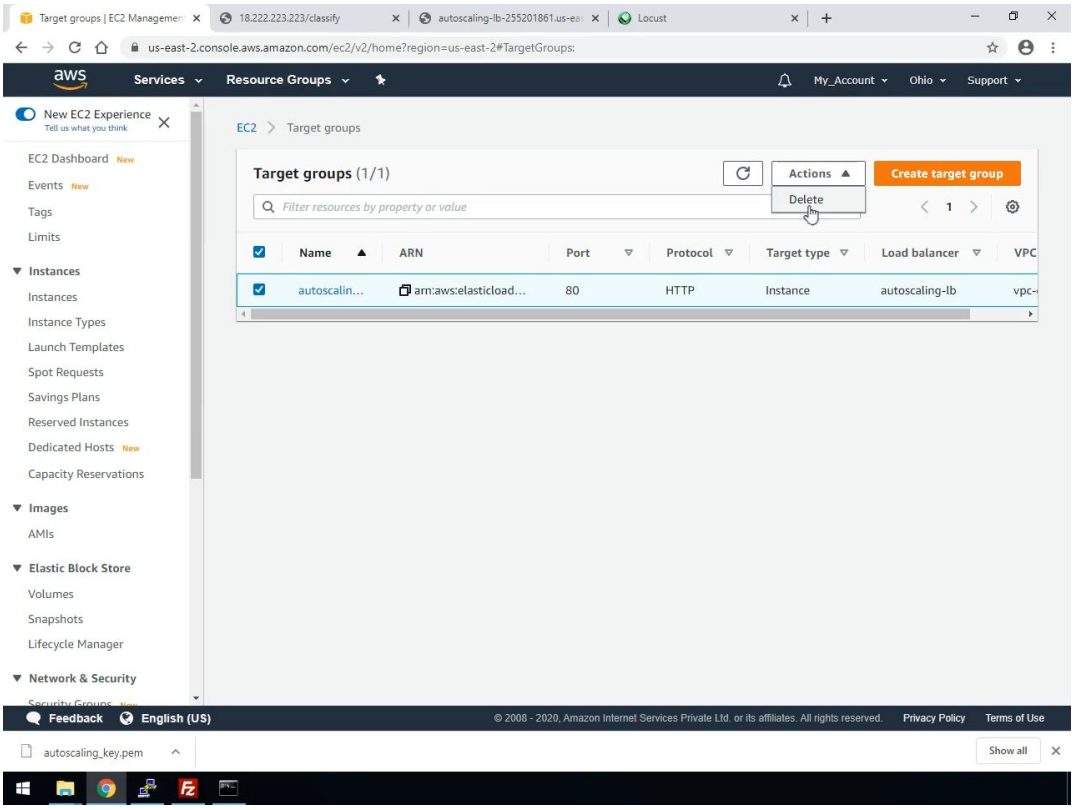


11:59

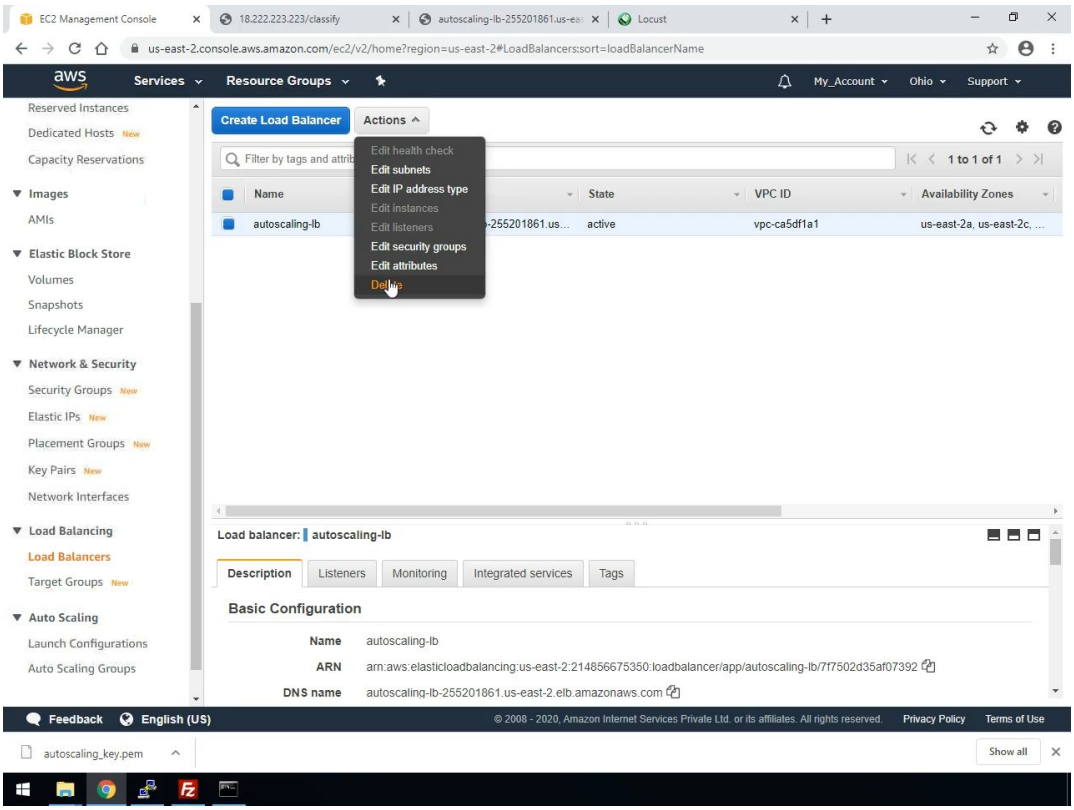
EC2 Management Console
18.222.223.223/classify
autoscaling-lb-255201861.us-ea
Locust
us-east-2.console.aws.amazon.com/ec2/v2/home?region=us-east-2#LoadBalancers:sort=loadBalancerName
Services
Resource Groups
My_Account
Ohio
Support
Reserved Instances
Dedicated Hosts
Capacity Reservations
Images
AMIs
Elastic Block Store
Volumes
Snapshots
Lifecycle Manager
Network & Security
Security Groups
Elastic IPs
Placement Groups
Key Pairs
Network Interfaces
Load Balancing
Load Balancers
Target Groups
Auto Scaling
Launch Configurations
Auto Scaling Groups
Create Load Balancer
Actions
Edit health check
Edit subnets
Edit IP address type
Edit instances
Edit listeners
Edit security groups
Edit attributes
Delete
Filter by tags and attrib
1 to 1 of 1
Name
State
VPC ID
Availability Zones
autoscaling-lb
active
vpc-ca5df1a1
us-east-2a, us-east-2c, ...
Load balancer: autoscaling-lb
Description
Listeners
Monitoring
Integrated services
Tags
Basic Configuration
Name
autoscaling-lb
ARN
arn:aws:elasticloadbalancing:us-east-2:214856675350:loadbalancer/app/autoscaling-lb/7f7502d35af07392
DNS name
autoscaling-lb-255201861.us-east-2.elb.amazonaws.com
Feedback
English (US)
© 2008 - 2020, Amazon Internet Services Private Ltd. or its affiliates. All rights reserved.
Privacy Policy
Terms of Use
autoscaling_key.pem
Show all

12:25

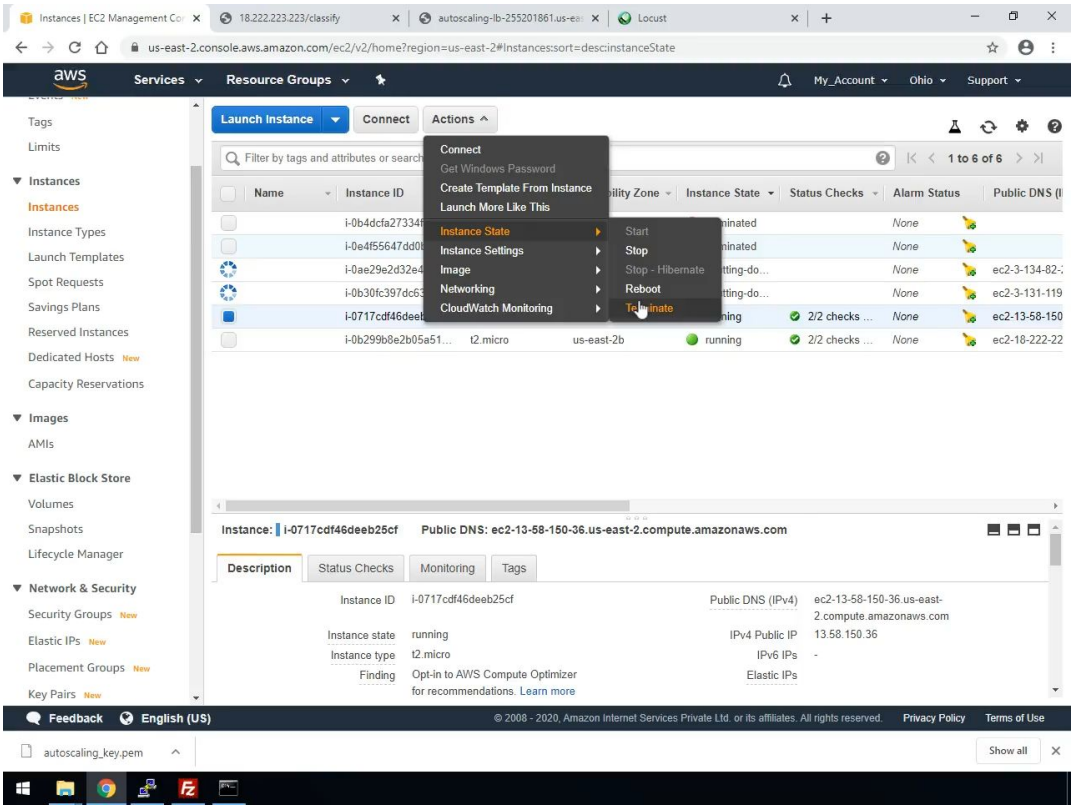
Instances | EC2 Management Co
18.222.223.223/classify
autoscaling-lb-255201861.us-ea
Locust
us-east-2.console.aws.amazon.com/ec2/v2/home?region=us-east-2#Instances:sort=desc:instanceState
Services
Resource Groups
My_Account
Ohio
Support
Tags
Limits
Instances
Instance Types
Launch Templates
Spot Requests
Savings Plans
Reserved Instances
Dedicated Hosts
Capacity Reservations
Images
AMIs
Elastic Block Store
Volumes
Snapshots
Lifecycle Manager
Network & Security
Security Groups
Elastic IPs
Placement Groups
Key Pairs
Launch Instance
Connect
Actions
Connect
Get Windows Password
Create Template From Instance
Launch More Like This
Instance State
Instance Settings
Image
Networking
CloudWatch Monitoring
Start
Stop
Stop - Hibernate
Reboot
Terminate
Filter by tags and attributes or search
1 to 6 of 6
Name
Instance ID
Instance State
Status Checks
Alarm Status
Public DNS (I
i-0717cdf46deeb25cf
t2.micro
us-east-2b
running
2/2 checks ...
None
ec2-13-58-150
i-0b299b8e2b05a51...
t2.micro
us-east-2b
running
2/2 checks ...
None
ec2-18-222-22
Instance: i-0717cdf46deeb25cf
Public DNS: ec2-13-58-150-36.us-east-2.compute.amazonaws.com
Description
Status Checks
Monitoring
Tags
Instance ID
i-0717cdf46deeb25cf
Instance state
running
Instance type
t2.micro
Finding
Opt-in to AWS Compute Optimizer for recommendations. Learn more
Public DNS (IPv4)
ec2-13-58-150-36.us-east-2.compute.amazonaws.com
IPv4 Public IP
13.58.150.36
IPv6 IPs
-
Elastic IPs
Feedback
English (US)
© 2008 - 2020, Amazon Internet Services Private Ltd. or its affiliates. All rights reserved.
Privacy Policy
Terms of Use
autoscaling_key.pem
Show all

13:18

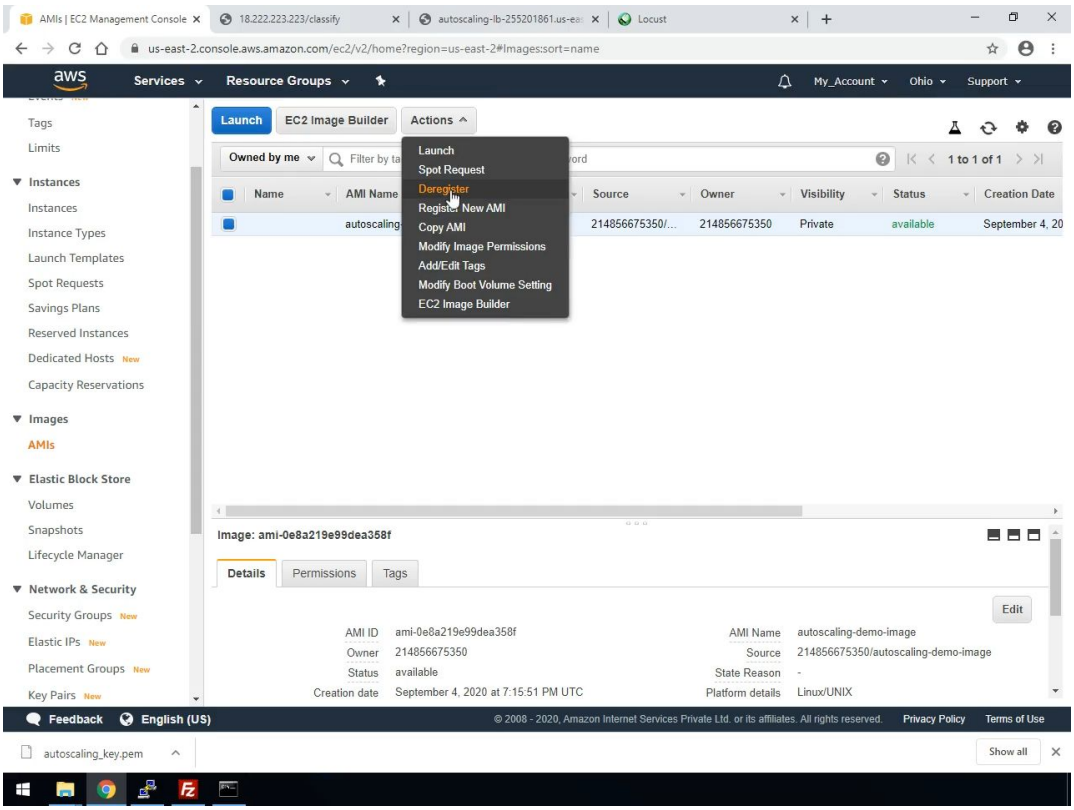
AMIs | EC2 Management Console
18.222.223.223/classify
autoscaling-lb-255201861.us-ea
Locust
us-east-2.console.aws.amazon.com/ec2/v2/home?region=us-east-2#Images:sort=name
Services
Resource Groups
My_Account
Ohio
Support
Launch
EC2 Image Builder
Actions
Launch
Spot Request
Deregister
Register New AMI
Copy AMI
Modify Image Permissions
Add/Edit Tags
Modify Boot Volume Setting
EC2 Image Builder
Owned by me
1 to 1 of 1
Name
AMI Name
Source
Owner
Visibility
Status
Creation Date
autoscaling
214856675350/...
214856675350
Private
available
September 4, 20
Image: ami-0e8a219e99dea358f
Details
Permissions
Tags
Edit
AMI ID
ami-0e8a219e99dea358f
Owner
214856675350
Status
available
Creation date
September 4, 2020 at 7:15:51 PM UTC
AMI Name
autoscaling-demo-image
Source
214856675350/autoscaling-demo-image
State Reason
-
Platform details
Linux/UNIX
Tags
Limits
Instances
Instance Types
Launch Templates
Spot Requests
Savings Plans
Reserved Instances
Dedicated Hosts
Capacity Reservations
Images
AMIs
Elastic Block Store
Volumes
Snapshots
Lifecycle Manager
Network & Security
Security Groups
Elastic IPs
Placement Groups
Key Pairs
Feedback
English (US)
© 2008 - 2020, Amazon Internet Services Private Ltd. or its affiliates. All rights reserved.
Privacy Policy
Terms of Use
autoscaling_key.pem
Show all

13:38

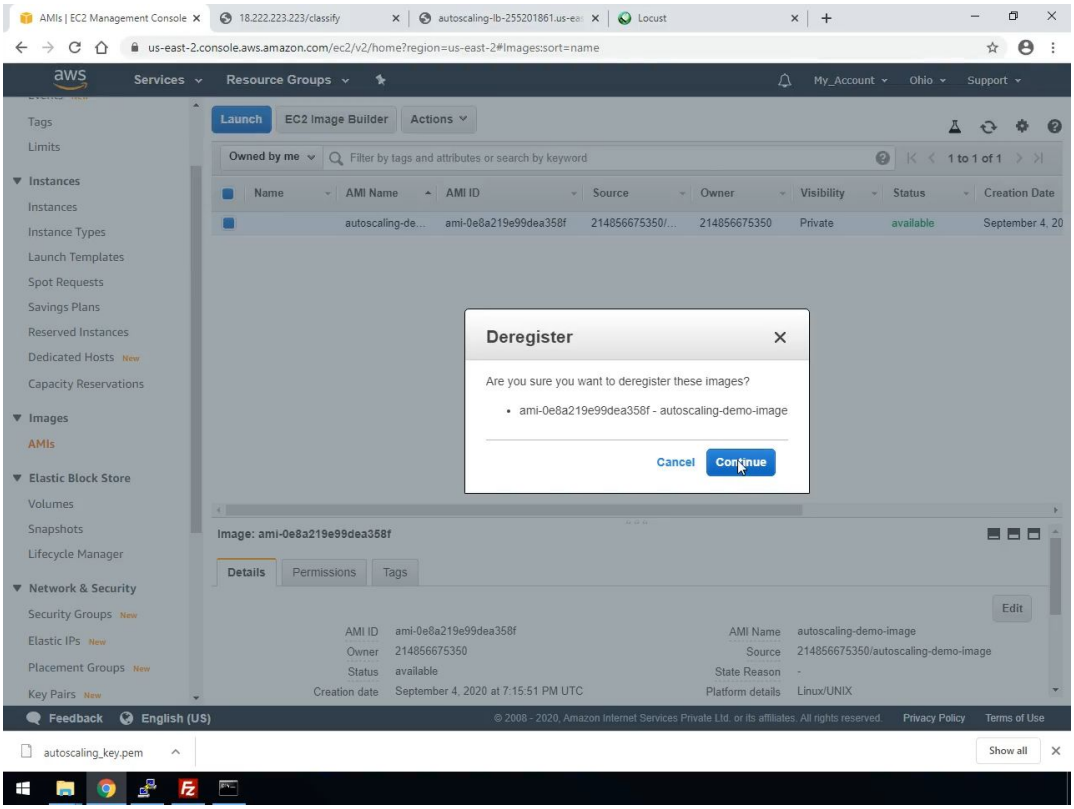
AMIs | EC2 Management Console
18.222.223.223/classify
autoscaling-lb-255201861.us-ea
Locust
us-east-2.console.aws.amazon.com/ec2/v2/home?region=us-east-2#Images:sort=name
Services
Resource Groups
My_Account
Ohio
Support
Launch
EC2 Image Builder
Actions
Owned by me
Filter by tags and attributes or search by keyword
1 to 1 of 1
Name
AMI Name
AMI ID
Source
Owner
Visibility
Status
Creation Date
autoscaling-de...
ami-0e8a219e99dea358f
214856675350/...
214856675350
Private
available
September 4, 20
Deregister
Are you sure you want to deregister these images?
ami-0e8a219e99dea358f - autoscaling-demo-image
Cancel
Continue
Image: ami-0e8a219e99dea358f
Details
Permissions
Tags
Edit
AMI ID
ami-0e8a219e99dea358f
Owner
214856675350
Status
available
Creation date
September 4, 2020 at 7:15:51 PM UTC
AMI Name
autoscaling-demo-image
Source
214856675350/autoscaling-demo-image
State Reason
-
Platform details
Linux/UNIX
Feedback
English (US)
© 2008 - 2020, Amazon Internet Services Private Ltd. or its affiliates. All rights reserved.
Privacy Policy
Terms of Use
autoscaling_key.pem
Show all

13:42